

How Texel Density Is Lowered with Mipmaps and LOD Bias

Rameel Malik

August 12, 2025

4-Step Explainer

1. Before rendering: build the mipmap pyramid.

After uploading the base texture (e.g., 4096×4096), calling `glGenerateMipmap(GL_TEXTURE_2D)` creates a chain of downscaled images stored in VRAM:

Level 0: 4096^2 , Level 1: 2048^2 , Level 2: 1024^2 , Level 3: 512^2 , ..., Level n : 1^2 .

These are precomputed once and reused every frame.

2. During rendering: the GPU chooses a mip level automatically.

For each fragment (pixel), the texture sampler estimates the on-screen footprint from UV derivatives and picks a mip level often summarized by

$$\lambda \approx \log_2(\rho),$$

where ρ encodes how fast the UVs change (times texture size). Intuition: large on-screen \Rightarrow small λ (higher-res level); small on-screen \Rightarrow large λ (lower-res level).

3. Why it looks blurrier at distance (and why that's good).

Lower mip levels are prefiltered averages of larger texel neighborhoods (Level 1 $\approx 2 \times 2$, Level 2 $\approx 4 \times 4$, etc.). Sampling them reduces shimmering/aliasing and lowers memory bandwidth, because fewer/smaller texels are fetched per screen pixel.

4. LOD bias: nudging the chosen level for performance or sharpness.

We apply a user-controlled offset to the hardware's choice:

$$\lambda_{\text{effective}} = \lambda_{\text{hardware}} + \text{bias}.$$

In code: `glSamplerParameterf(samp, GL_TEXTURE_LOD_BIAS, bias)`. A **positive** bias (e.g., `+1.0`) picks lower-resolution mips sooner (\approx one level lower, half-res each dimension) \Rightarrow *lower texel density, higher performance*. A small **negative** bias (e.g., `-0.25`) keeps higher-res mips a bit longer \Rightarrow *sharper but costlier*.